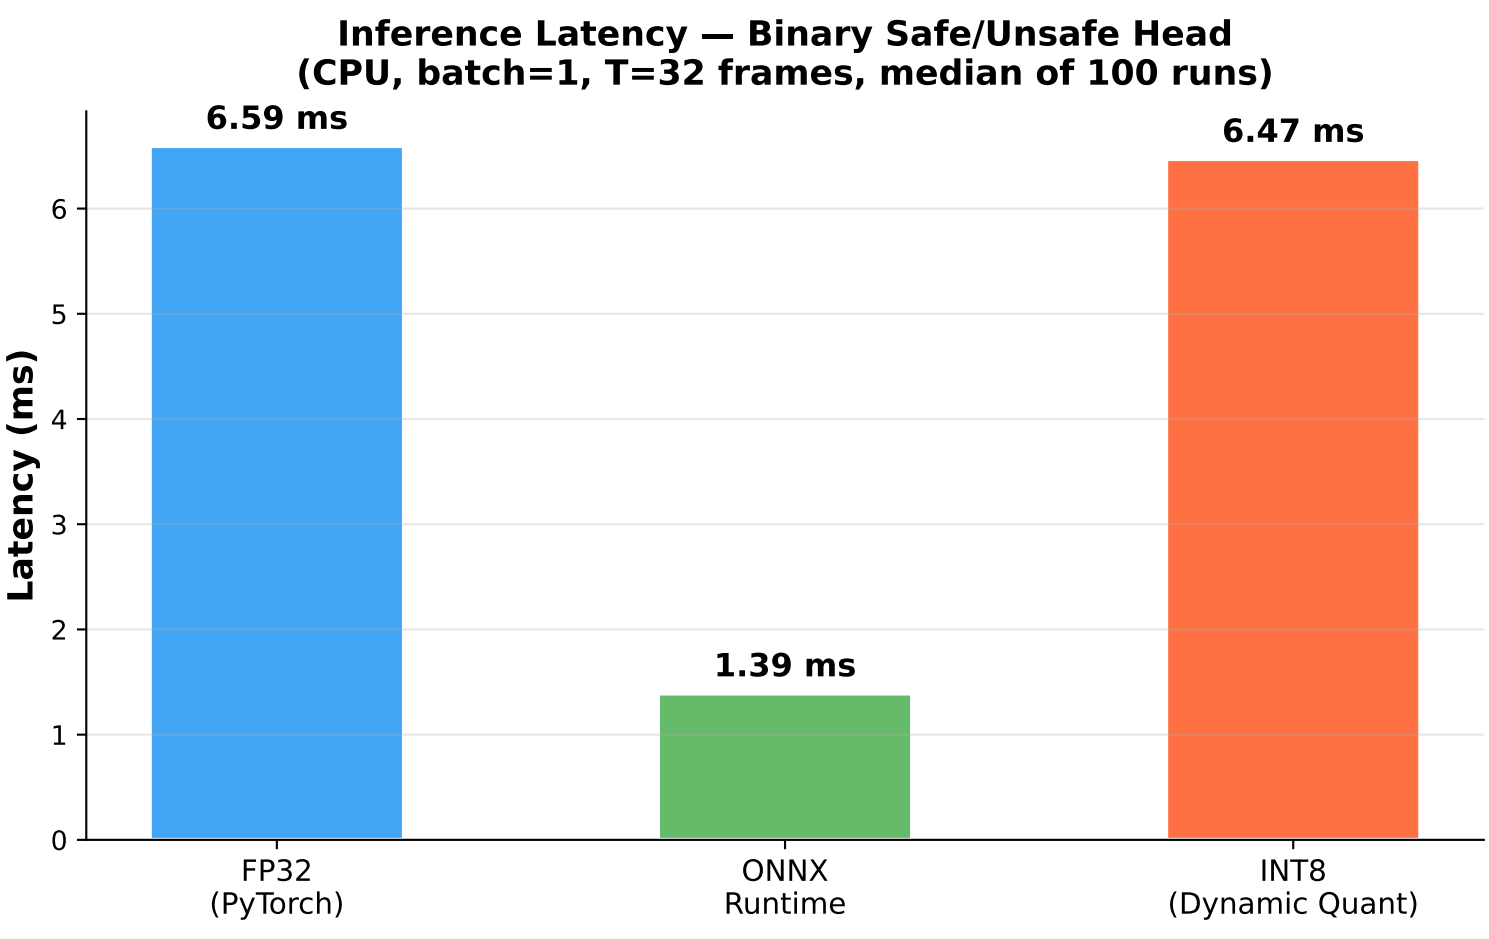
Inference Latency — Binary Safe/Unsafe Head
(CPU, batch=1, T=32 frames, median of 100 runs)
Latency (ms)
0
1
2
3
4
5
6
6.59 ms
1.39 ms
6.47 ms
FP32
(PyTorch)
ONNX
Runtime
INT8
(Dynamic Quant)